



# Building a personal search engine with llama-index

Judith van Stegeren & Yorick van Pelt

# Before this tutorial

1. Download the files from <https://github.com/datakami/pydata-llama-index-tutorial/>
2. Follow the readme to get the requirements
  - a. Local Python environment
  - b. Docker
  - c. Google Colab

# Tutorial (1 hour)

## **Part 1: searching with llama-index**

Notebook 1: retrieval and querying

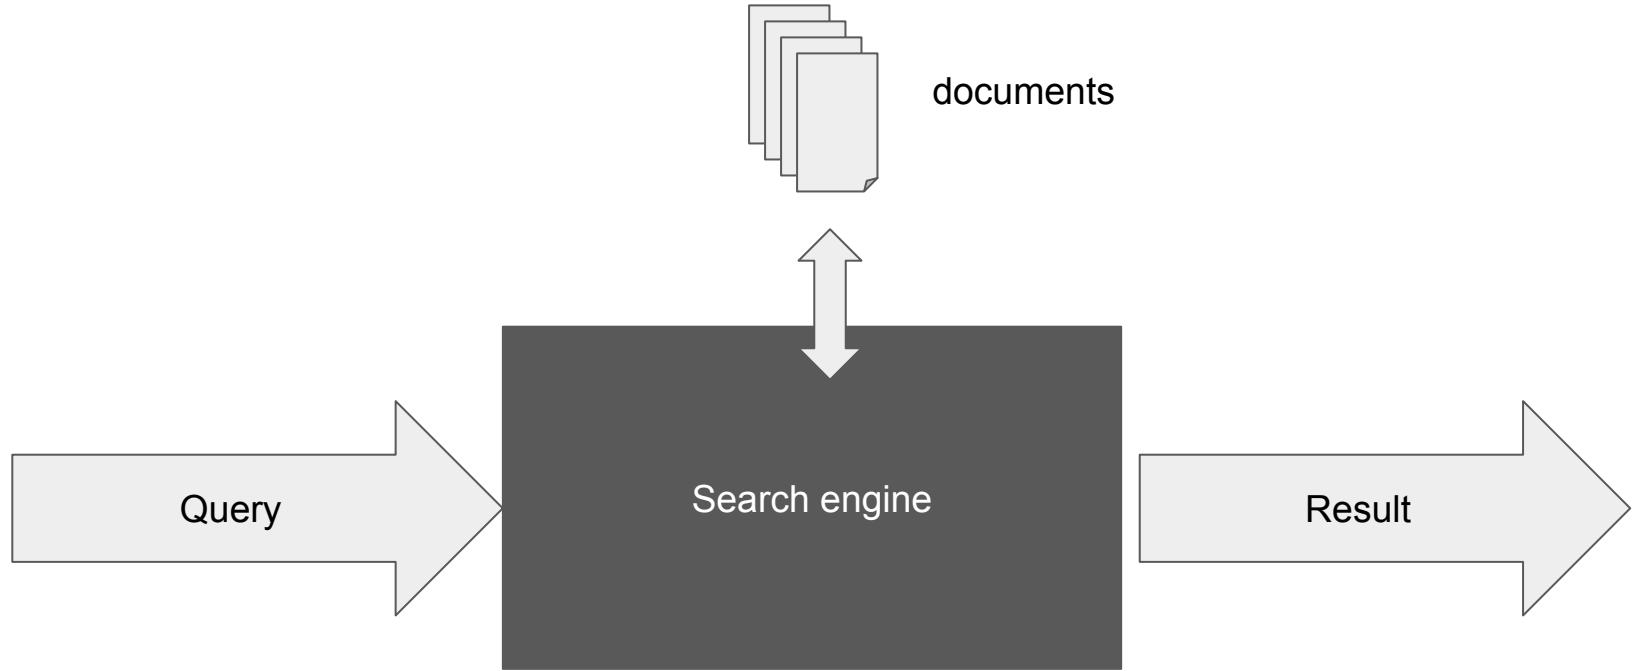
## **Part 2: building a search engine from your own documents**

Notebook 2: building an index from text documents

## **Wrap-up: things to explore after this tutorial**

# Part 1: searching with llama-index

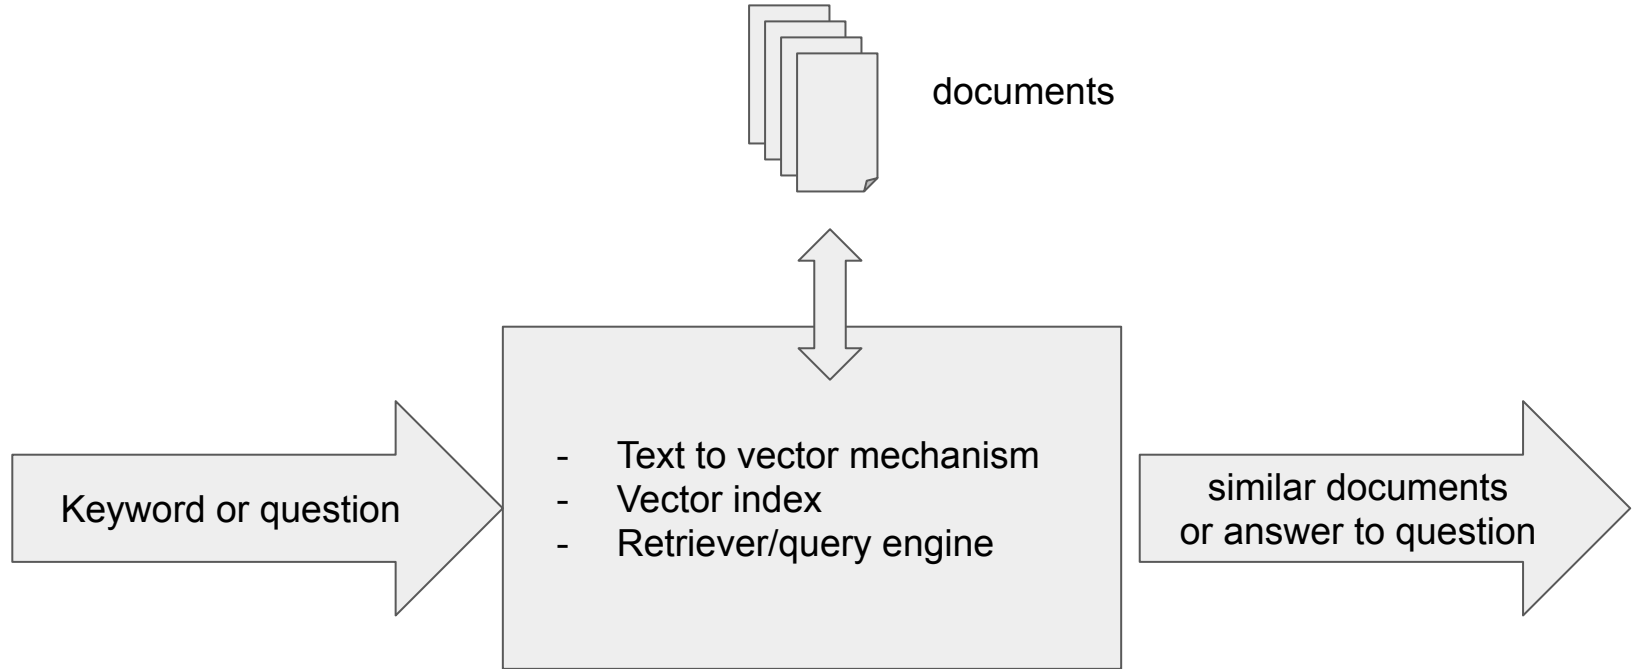
# Personal search engine



```
import llama_index
```



# Personal search engine



# What is this “vector store index” thing?

Vector: row of numbers

Vector store: *thing* that stores vectors

Index: a list of items, each of which identifies a particular record in a computer file or database and contains information about its address



# Retrieving vs querying

## Retrieval

Ingredients: index, query

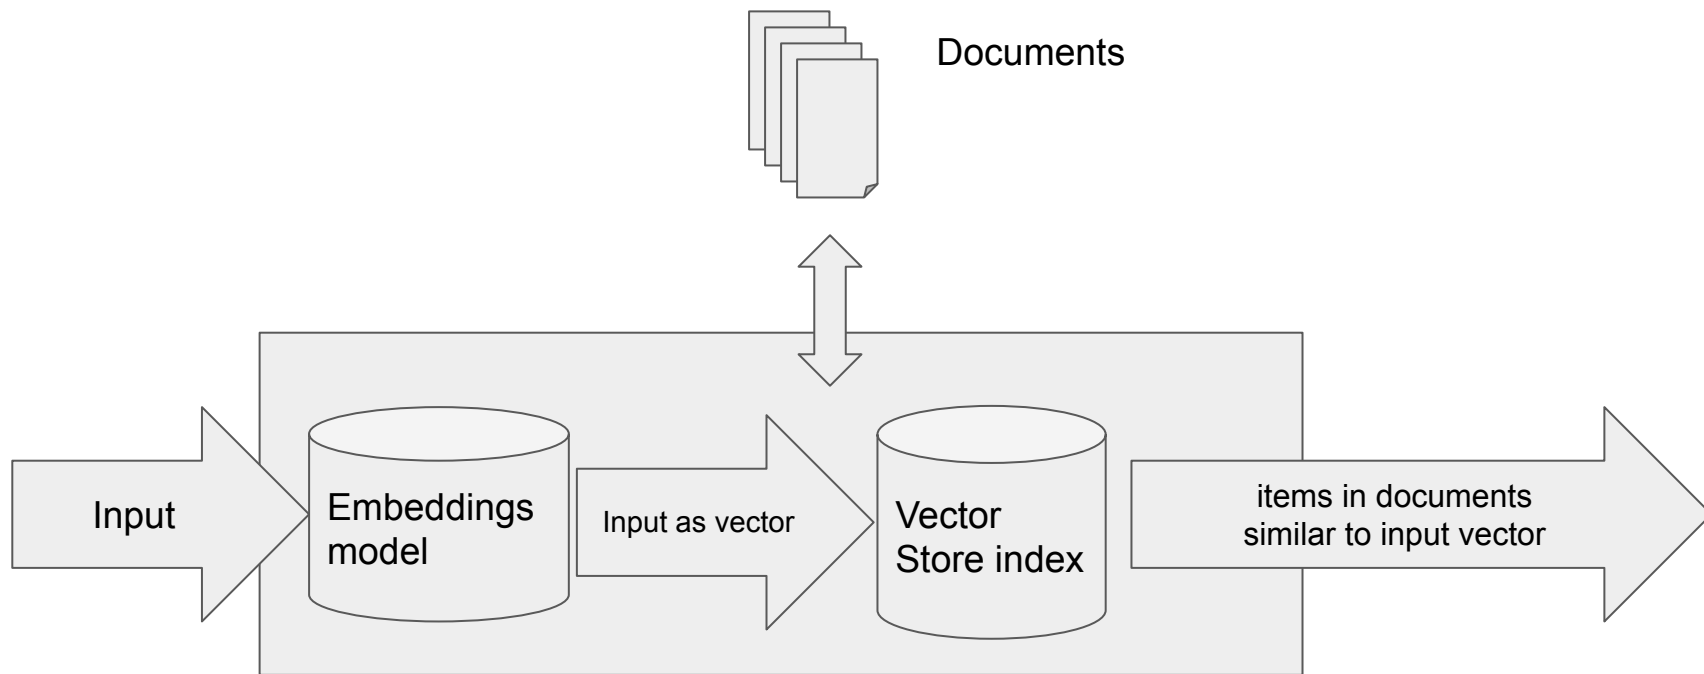
Result: most "similar" things in the index

## Querying

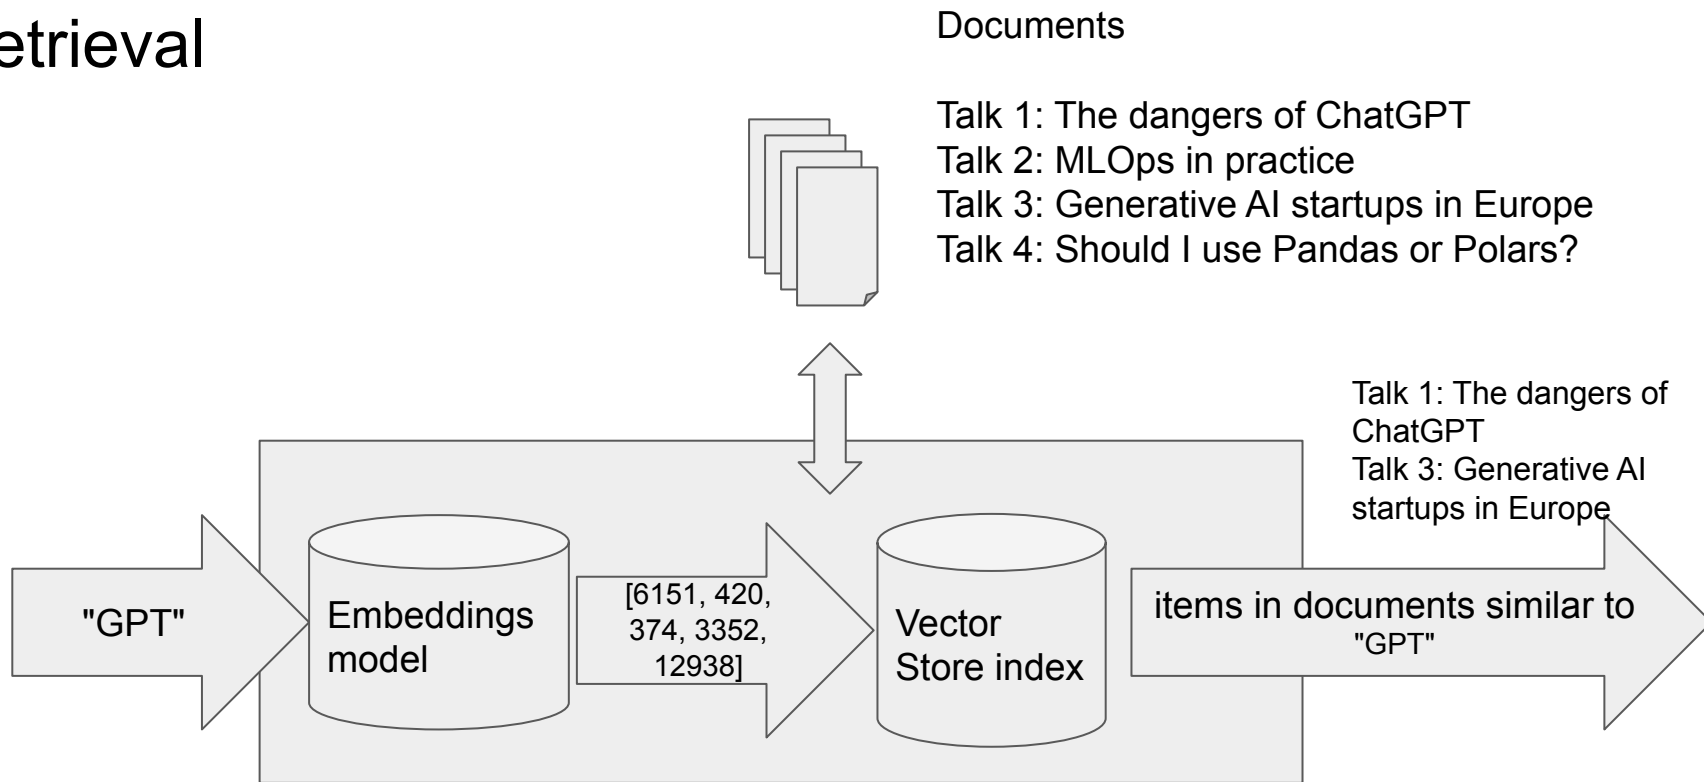
Ingredients: index, query, context

Result: LLM uses the context + retrieval result to write an answer to the query


# Retrieval



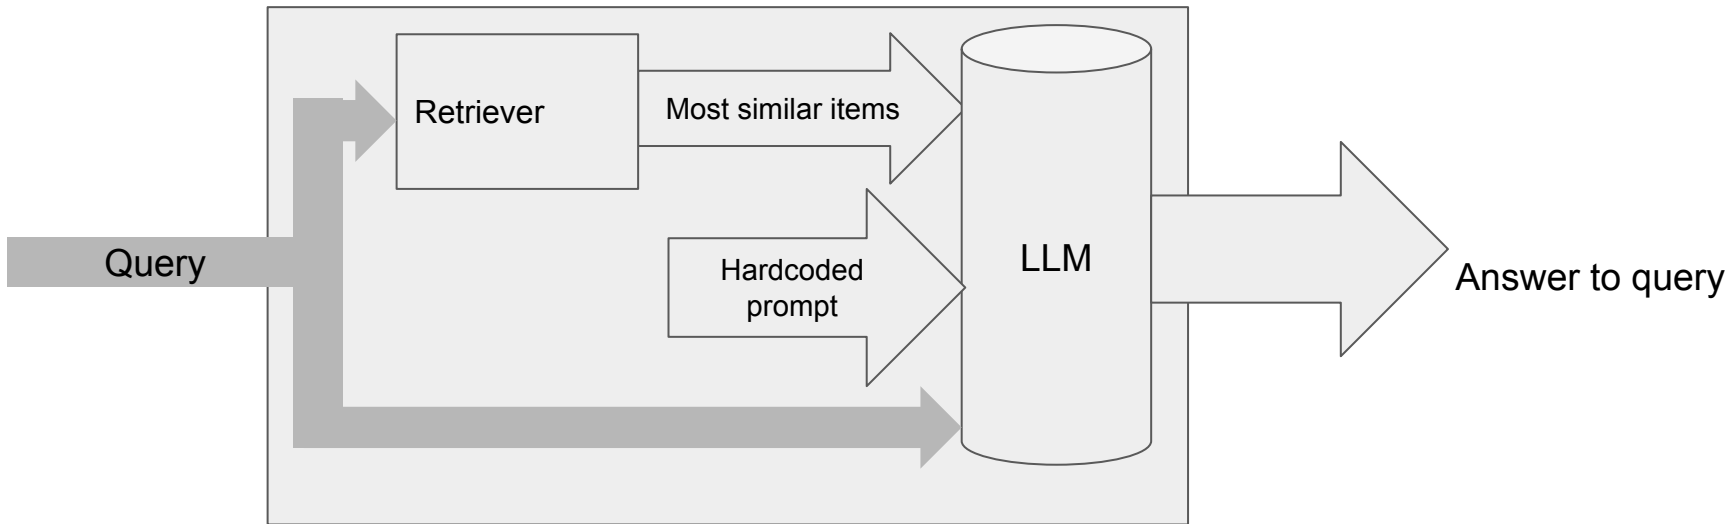

# Retrieval




# Querying



Documents



# Querying



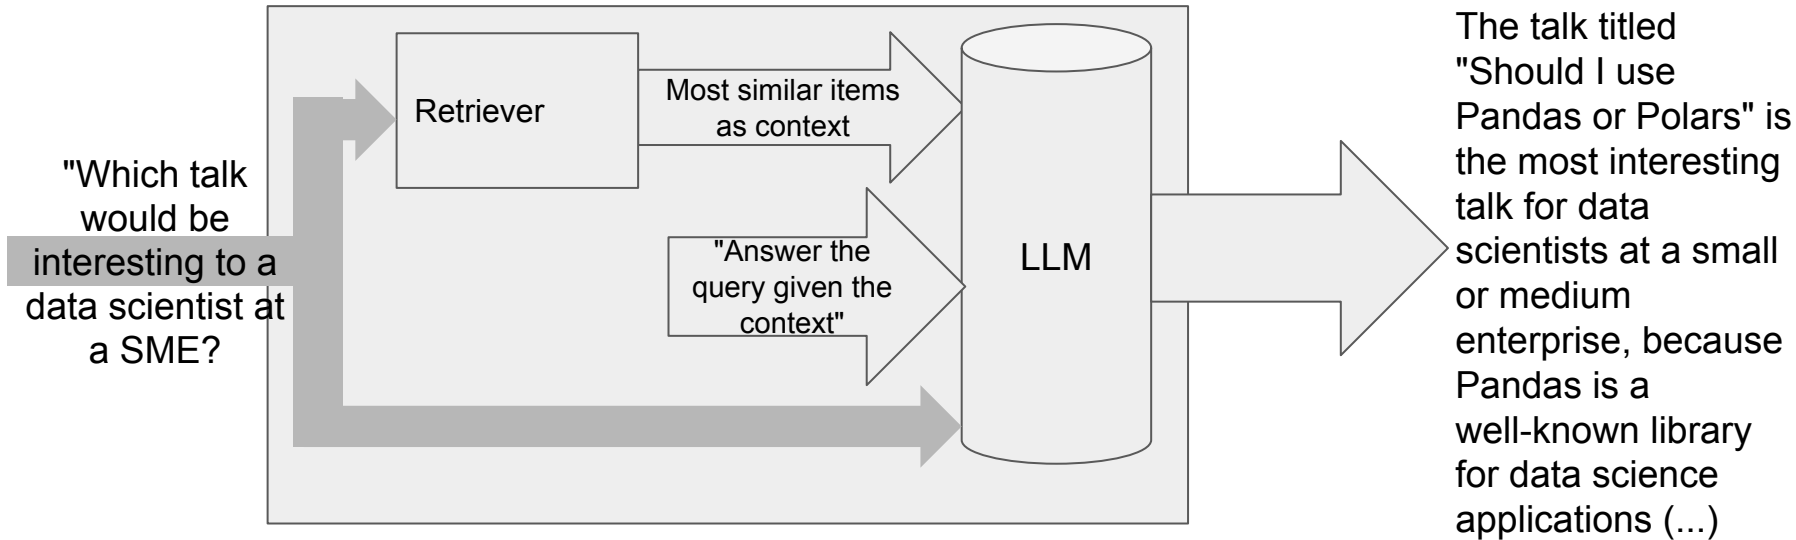
## Documents

Talk 1: The dangers of ChatGPT

Talk 2: MLOps in practice

Talk 3: Generative AI startups in Europe

Talk 4: Should I use Pandas or Polars?



# Our first index

`index` is a vector store index with PyData Amsterdam 2023 talks!

Index consists of document snippets called `nodes`.

# Our first index

A node has text and metadata

```
node.text
```

```
"Building a personal search engine with llama-index\n\nWouldn't it be great to have a Google-like search engine,  
but then for your own text files and completely private?  
In this tutorial (...)"
```

```
node.metadata
```

```
{'title': 'Building a personal search engine with llama-index',  
'speakers': 'Judith van Stegeren, Yorick van Pelt'}
```

# Retrieving with llama-index

`index` is a vector store index with the PyData schedule

```
retriever = index.as_retriever()
```

```
search_results = retriever.retrieve("your query here")
```



# Dealing with search results

`index` is a vector store index with the PyData schedule

`search_results` is a list of retrieval results

A `result` has a `node` and a `score`

```
first_result = search_results[0]
```

```
first_result.node
```

```
first_result.score
```

# Querying with llama-index

`index` is a vector store index with the PyData schedule

```
query_engine = index.as_query_engine()
```

```
search_results = query_engine.query("your query here")
```

# Time for a notebook: **Exercises-1.ipynb**

`index` is a vector store index with the PyData schedule

```
node.text
```

```
node.metadata
```

```
retriever = index.as_retriever()
```

```
search_results = retriever.retrieve("your query here")
```

```
query_engine = index.as_query_engine()
```

```
search_results = query_engine.query("your query here")
```

```
result.node
```

```
result.score
```

Tutorial repository: <https://github.com/datakami/pydata-llama-index-tutorial>

## Part 2: building a vector index

# Recap: our first index

`index` is a vector store index with PyData Amsterdam 2023 talks!

Index consists of document snippets called nodes.

A `node` has `text` and `metadata`

# Building a new vector store index

An index is created from a set of Documents.

Documents are split into Nodes (we've seen these before!)

Four steps:

1. Load data from source (vanilla Python)
2. Transform data to "Document" format
3. Put documents in vector store index
4. Save the vector store index to file

# Step 1: loading data

from plaintext file, from json, from the web, etc.

All kinds of custom data loaders on <https://llamahub.ai/>. (Quality varies)

**# plaintext**

```
with open("my_research_notes.txt", 'r') as infile:
    research_notes = infile.read()
```

**# json**

```
with open("movie_descriptions.json", 'r') as infile:
    movies = json.loads(infile.read())
```

**# from web**

```
response = requests.get("httpx://myfavoriteapi.com/endpoint?query=somedata")
api_data = json.loads(response.text)
```

# Step 2: creating Documents from data

Transform data to a structure that llama\_index understands.

All Document properties are completely optional!

```
document = llama_index.Document() # empty document
```

```
Document(  
    id_='82c3f6b8-4e64-4d76-9629-f51aaa7446a0',  
    embedding=None,  
    metadata={},  
    excluded_embed_metadata_keys=[],  
    excluded_llm_metadata_keys=[],  
    relationships={},  
    hash='44136fa355b3678a1146ad16f7e8649e94fb4fc21fe77e8310c060f61caaff8a',  
    text='',  
    start_char_idx=None,  
    end_char_idx=None,  
    text_template='{metadata_str}\n\n{content}',  
    metadata_template='{key}: {value}',  
    metadata_seperator='\n'  
)
```



## Step 2: creating Documents from data

Transform data to a structure that llama\_index understands.

All Document properties are completely optional!

```
document = llama_index.Document() # empty document
```

```
Document(  
    id_='82c3f6b8-4e64-4d76-9629-f51aaa7446a0',  
    metadata={},  
    text='',  
)
```

## Step 2: creating Documents from data

Simplest way to create a document: just add text.

```
simple_text = "Hello world"

document = llama_index.Document(text=simple_text)

Document(
    id_='993d3bfd-83df-4fb5-b6f7-5ea688df8032',
    metadata={},
    text='Hello world',
)
```

## Step 2: creating Documents from data

You can add custom identifiers to Documents.

This is good for tracking Document sources!

Important: use `id_` instead of `id` (this will fail silently)

```
my_filename = "somefile.txt"
with open(my_filename, "r") as infile:
    my_text = infile.read()

document = llama_index.Document(text=my_text, id_=my_filename)

Document(
    id_="somefile.txt",
    metadata={},
    text="The contents of a file on my computer",
)
```

# Step 2: creating Documents from data

## Add metadata as dictionary

```
simple_text = "Hello world"
some_metadata = { "author" : "Judith", "created_at": datetime.now() }

document = llama_index.Document(text=simple_text, metadata=some_metadata)

Document(
  id_='70b4c14c-18e9-49c8-9265-01760b8d5c3d',
  metadata={
    'author': 'Judith',
    'created_at': datetime(2023, 9, 8, 11, 7, 55, 281502)
  },
  text='Hello world'
)
```

## Step 3: creating vector index from data

```
documents = [Document(text=datapoint) for datapoint in data]
```

```
index = VectorStoreIndex.from_documents(documents)
```

And we have an index!

Now we can create a `retriever` or `query_engine` and do everything we saw in part 1.

## Step 3: creating vector index from data

We can use a `ServiceContext` to change:

- the size of text in nodes, or chunk size
- the overlap between nodes
- the embeddings model for transforming text to numbers
- the language model used for querying

And lots of other things as well!

## Step 3: creating vector index from data

```
my_service_context = ServiceContext.from_defaults()
my_documents = [Document(text=datapoint) for datapoint in data]

index = VectorStoreIndex.from_documents(documents=my_documents,
service_context=my_service_context)
```

In part 1 of this tutorial, we used a ServiceContext to

- use the smallest sentence-transformers embeddings model
- remove the default LLM so llama-index won't call OpenAI's GPT

## Step 3: creating vector index from data

```
service_context = ServiceContext.from_defaults()

service_context = ServiceContext.from_defaults(llm=None)

# llama-index will download the model for us
my_model = "local:sentence-transformers/all-minilm-l6-v2"
service_context = ServiceContext.from_defaults(embed_model=my_model)

service_context = ServiceContext.from_defaults(chunk_size=10)
```



## Step 4: saving the index to file

`StorageContext` is a utility class for storing indices, vectors, and nodes.

`index` is a `VectorStoreIndex`

```
index_path = 'indices/name_of_my_index'  
index.storage_context.persist(index_path)
```

## Step 4: saving the index to file

We can load an index from disk by recreating its `StorageContext` and `ServiceContext`.

```
index_path = 'indices/my_old_index'

storage_context = StorageContext.from_defaults(persist_dir=index_path)
service_context = ServiceContext.from_defaults(llm=my_llm)
my_old_index = load_index_from_storage(storage_context,
service_context=service_context)
```

# Time for a notebook: **Exercises-2.ipynb**

`index` is a vector store index

```
node.text
node.metadata
```

```
index.as_retriever()
retriever.retrieve("your query here")
```

```
index.as_query_engine()
query_engine.query("your query here")
```

```
result.node
Result.score
```

```
Document(
    text="hello",
    id_="hello.txt",
    metadata={"author": "Judith"}
)
```

```
from llama_index.llms import OpenAI
my_llm = OpenAI(model="gpt-4")
```

```
my_embed_model =
"local:sentence-transformers/all-minilm-l6-v2"
```

```
ServiceContext.from_defaults(
    llm=my_llm, embed_model=my_embed_model,
    chunk_size=my_chunk_size
)
```

```
VectorStoreIndex.from_documents(
    documents=documents,
    service_context=my_service_context
)
```


```
index.storage_context.persist('my_path')
```

```
StorageContext.from_defaults(persist_dir=index_path)
```

```
load_index_from_storage(
    storage_context,
    service_context=service_context
)
```

Tutorial repository: <https://github.com/datakami/pydata-llama-index-tutorial>

# What's next?



**datakami**  
research & development

Experiment with your own data and favorite models

Other stuff to change or optimize: chunking, other models, meta-data, hardcoded prompts, ...

Use cases: research notes, diaries, CRM, improve other LLM pipelines, ...

Please don't use this for production

Defaults are slowly getting more sane over time

Want to understand *embeddings* better? [https://vickiboykis.com/what\\_are\\_embeddings/](https://vickiboykis.com/what_are_embeddings/)

Tutorial repository: <https://github.com/datakami/pydata-llama-index-tutorial>

